

Assignment 1: Implementing and Analyzing a Basic RL Algorithm

Due: Sunday May 26, before midnight EST

Objective:

To gain hands-on experience in implementing a basic Markov Decision Process (MDP) and applying a Dynamic Programming algorithm, such as Policy Iteration or Value Iteration.

Tasks:

- I. Implement a Gridworld Environment: **(20 points)**
 - a. Create a gridworld environment, a simple and commonly used model in RL.
 - b. Define states, actions, transition probabilities, and rewards.
 - c. Include varied rewards, with some states having negative rewards (obstacles) and others positive rewards (goals).
- II. A linear solver to solve the system $Ax = b$ for the deterministic case. **(10 points)**
- III. Apply a Dynamic Programming Algorithm: **(30 points)**
 - a. Implement Policy Iteration or Value Iteration to find an optimal policy for navigating the gridworld.
 - b. Ensure your implementation accurately reflects the theoretical principles of the chosen algorithm.
- IV. Analyze Algorithm Convergence: **(10 points)**
 - a. Examine the number of iterations your algorithm takes to converge.
 - b. Discuss factors affecting convergence speed and policy quality.

Submission:

Python Code: Fully documented code implementing the gridworld and the Dynamic Programming algorithm. **(The points divided in the tasks total of 70 points)**

Report: A comprehensive report detailing the implementation, the algorithm's results, and an analysis of its convergence and effectiveness. **(30 points)**

The Problem Description:

Consider the following Gridworld example. There are 4 different actions (north, south, east, west). for all the states (cells in the grid), each one of the actions (north, south, east, and west) is chosen with probability $\frac{1}{4}$. The agent then moves with probability 1 to the chosen direction. While moving, if the agent hits a wall to go off-grid or the black cells (blocked), it cannot move and it receives a reward of 0; if moving to a cell in the grid, the reward is 0.; if it reaches red

cells (fire), the game will terminate and receives a penalty -5; and, if it reaches the goal cell (G) it receives a reward +5 and the game will terminate.

State Space: Each cell in the grid represents a state. The agent can be in any cell except for those marked as blocked.


Actions: The agent can move up, right, down, or left. Moving into a blocked cell is not allowed, and attempting to move outside the grid bounds results in staying in the same cell.

Transition Model: The agent moves in the chosen direction with a probability of “1 - noise”. With probability noise, it may end up moving in a different direction.

Rewards: The agent receives a specific reward (positive or negative) when it enters a cell. Moving into a danger cell yields a negative reward, reaching the goal cell yields a positive reward, and other moves typically have no reward (zero).

Termination: The episode ends if the agent reaches the goal cell or a danger cell.

Note: You are given a code snippet. Feel free to use it or use your own implementation.



You are required to write 2 programs (in Python 3) to find the state-value for each one of the states for discount rates of 0.95 and 0.75. You are also required to find the state-value for each state when the grid is noisy (probability of resulting state not being what was expected).

The 2 programs you will write are:

1. A linear solver to solve the system $Ax = b$ for only deterministic case.
2. A dynamic programming approach (policy iteration or value iteration).